

## SPMV Task 3: Performance Analysis

### Introduction

To prepare the code for analysis, I removed functions which were used with ICC (Intel C Compiler) in Task 2, as they were not required for this task. The primary objective of this work is to analyze the performance of Sparse Matrix-Vector Multiplication (SpMV) with a focus on two key aspects:

#### 1. Vectorization Analysis

- o Determine if the relevant parts of the routines are being autovectorized by the compiler.
- o Identify reasons for lack of autovectorization, if applicable.
- o Explore ways to assist the compiler in achieving autovectorization.

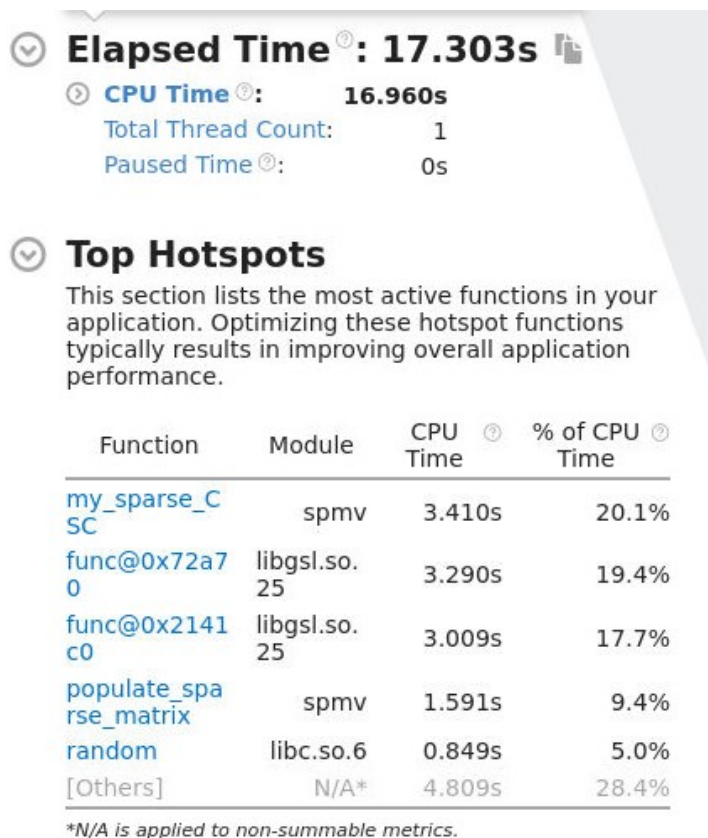
#### 2. Memory and Cache Behavior Analysis

- o Assess whether heap usage is being managed efficiently (e.g., avoiding memory leaks and ensuring proper allocation sizes).
- o Analyze the memory access pattern in the code to understand its impact on cache performance, focusing on spatial and temporal locality.

The analyses will be conducted using Intel VTune Profiler.

### Hotspots Analysis

As the first step, I performed a **hotspots analysis** of my code to identify the most computationally intensive sections. This serves as a baseline for further optimization and performance tuning.



The function that consumes the most CPU time is `my_sparse_csc`, followed by two functions from a library. Since the library functions are external, there is nothing we can optimize there.

For the `populate_sparse_matrix` function, its purpose is to populate the sparse matrix, and there is nothing significant to optimize in this routine.

<code>for (unsigned int i = 0; i &lt; n * n; i++) {</code>	0.4%	59.99
<code>if ((rand() % 100) / 100.0 &lt; density) {</code>	5.5%	929.24
<code>// Get a pseudorandom value between -9.99 e 9.99</code>		

Additionally, `random`, which is also from a library, cannot be optimized either.

Let's focus on the `my_sparse_csc` function.

<code>// Iterate over columns for CSC format</code>		
<code>for (unsigned int j = 0; j &lt; n; j++) {</code>		
<code>  sparse[j].col = k; // Column pointer: Start of each column in `val` and `row`</code>		
<code>  for (unsigned int i = 0; i &lt; n; i++) {</code>	0.6%	
<code>    if (mat[i * n + j] != 0) {</code>	17.6%	
<code>      sparse[k].row = i; // Store the row index for the non-zero element</code>	1.6%	
<code>      sparse[k].val = mat[i * n + j]; // Store the non-zero element itself</code>	0.2%	
<code>      k++;</code>	0.1%	
<code>    }</code>		
<code>  }</code>		
<code>}</code>		

It seems to consume a lot of time due to the `mat[i * n + j]` operation, where we traverse the matrix column by column instead of row by row. In C, matrices are stored in row-major order (like a flattened table), so accessing elements row by row would typically be more efficient. Unfortunately, since this function is responsible for creating the CSC (Compressed Sparse Column) format matrix, column-wise traversal is required, and we cannot change this access pattern without altering the matrix format.

## Memory Consumption

Memory Consumption ⓘ ⓘ

Analysis Configuration

Collection Log

Summary

Bottom-up

INTEL VTUNE P

⌵ Elapsed Time ⓘ: 17.616s

Allocation Size:

Deallocation Size:

Allocations:

Total Thread Count:

Paused Time ⓘ:

18.2 GB

2.5 GB

68

1

0s

⌵ Top Memory-Consuming Functions

This section lists the most memory-consuming functions in your application.

Function	Memory Consumption	Allocation/Deallocation Delta	Allocations	Module
main	15.0 GB	12.9 GB	7	spmv
func@0x2bc1e0	1.2 GB	1.2 GB	9	libopenblas.so.0
gsl_spmatrix_alloc_nzmax	1.1 GB	751.7 MB	14	libgsl.so.25
func@0x214a50	859.0 MB	859.0 MB	1	libgsl.so.25
gsl_block_alloc	262.1 KB	0.0 B	2	libgsl.so.25
[Others]	83.9 KB	18.2 KB	35	N/A*

\*N/A is applied to non-summable metrics.

In the column "Allocation/Deallocation Delta," we can see that a significant amount of memory is not being deallocated. This indicates that we need to review the code to ensure proper deallocation of these data.

⌵ Top Memory-Consuming Functions

This section lists the most memory-consuming functions in your application.

Function	Memory Consumption	Allocation/Deallocation Delta	Allocations	Module
main	15.0 GB	0.0 B	7	spmv
func@0x2bc1e0	1.2 GB	1.2 GB	9	libopenblas.so.0
gsl_spmatrix_alloc_nzmax	1.1 GB	0.0 B	14	libgsl.so.25
func@0x214a50	859.0 MB	0.0 B	1	libgsl.so.25
gsl_block_alloc	262.1 KB	0.0 B	2	libgsl.so.25
[Others]	83.9 KB	18.2 KB	35	N/A*

\*N/A is applied to non-summable metrics.

After improving the code, almost all the issues were resolved. The missing `free()` function calls were added, addressing most of the memory leaks. However, there is still one function, `func@0x2bc1e0`, which does not deallocate memory. Unfortunately, I don't know which function this refers to since its name doesn't refer directly to it.

### Vectorization Analysis

HPC Performance Characterization

INTEL VTUNE PROFILER

Analysis Configuration

Collection Log

Summary

Bottom-up

Cache Bound: 8.3% of Clockticks

DRAM Bound: 14.1% of Clockticks

NUMA: % of Remote Accesses: 0.0%

Vectorization: 5.9% of Packed FP Operations

- Instruction Mix:
- SP FLOPs: 0.0% of uOps
  - DP FLOPs: 1.9% of uOps
  - x87 FLOPs: 0.0% of uOps
  - Non-FP: 98.1% of uOps
- FP Arith/Mem Rd Instr. Ratio: 0.070
- FP Arith/Mem Wr Instr. Ratio: 0.185

Top Loops/Functions with FPU Usage by CPU Time

This section provides information for the most time consuming loops/functions with floating point operations.

Function	CPU Time	% of FP Ops	FP Ops: Packed	FP Ops: Scalar	Vector Instruction Set	Loop Type
random	0.890s	4.5%	0.0%	100.0%		
[Loop at line 8 in my_dense]	0.400s	28.0%	0.0%	100.0%		
[Loop@0x3087b8 in func@0x308760]	0.130s	50.0%	100.0%	0.0%	AVX(256); FMA(256)	
[Loop at line 16 in my_coo]	0.090s	33.3%	0.0%	100.0%		
[Loop@0x22f100 in gsl_splbas_dgemv]	0.085s	20.0%	0.0%	100.0%		
[Others]	0.120s	27.3%	0.0%	100.0%		

\*N/A is applied to non-summable metrics.

The vectorization efficiency is reported to be only 5.9%.

Despite using the compiler flags `-g -c -Wall -Wextra -Ofast -ftree-vectorize`, the low vectorization rate persists. The reason for this behavior remains unclear and requires further investigation. This issue was already observed in Task 2, where different optimization flags were tested to evaluate whether vectorization improved execution speed. However, with GCC, no significant change in execution time was noted.